

Applied AI Workshop

Welcome!

The Team



Tarek Aziz, Ph.D.
Associate Professor, CCEE



Fred Livingston, Ph.D.
Associate Teaching
Professor, ISE



Edgar Lobaton, Ph.D.
Professor, ECE



Feiya Xiang



Geet Khatri



Andrew Chen

The Technical Team



Siqing Wei, Ph.D.
Assistant Professor
Engineering Education
Youngstown State Univ.

Yunmeng Han
University of Cincinnati



Ren Butler
Carnegie Mellon
University

Samson Quaye
Illinois Institute of
Technology



Technical Support

- Cong Gian (IU - JetStream)
- Julian Pistorius (IU - Jetstream)
- Shawn Dunning, Director of IT
- Andrew Petersen, OIT

Leadership & Support



Jim Pfaendtner, Ph.D.
Executive Vice Chancellor and Provost



Veena Misra, Ph.D.
Interim Dean of the College of Engineering

- Marissa Langford, Assistant Dean of College Operations
- Jack Boden, Senior Event Producer
- Sarah Hester-Smith, University Program Specialist

NAIRR Pilot

National Artificial Intelligence
Research Resource Pilot



Workshop Overview

Goal: This workshop aims to build practical, cross-disciplinary AI literacy across the College of Engineering, progressing participants from foundational concepts to advanced methods in machine learning, deep learning, and physics-informed AI.

Module 1

AI and Software
Engineering Tools

Module 2

Foundations of Machine
Learning

Module 3

Foundations of Deep
Learning

Module 4

Physics and Deep
Learning

Today's Agenda

Time	Activity	Lead
1:00 pm - 2:00 pm	Introduction to AI	Edgar Lobaton
2:00 pm - 2:30 pm	Hands on: Exploring Code Assistants	Edgar Lobaton
2:30 pm - 3:00 pm	Presentation: Playing with AI - Curiosity, Mistakes, and Finding Real Co-Intelligence	Tarek Aziz
	Tarek Aziz is an Associate Professor in Civil, Construction, and Environmental Engineering. He is an environmental process engineer interested in nature-based solutions for water treatment and water quality. His work involves modeling, experimentation, and field work and in the past years, he's woven AI into all three. In his spare time Tarek loves spending time with family, traveling, reading and playing video games.	
3:00 pm - 3:15 pm	Coffee Break	
3:15 pm - 3:45 pm	Tool Overview: Containers	Edgar Lobaton
3:45 pm - 4:15 pm	Tool Overview: GitHub	Tarek Aziz
4:15 pm - 5:00 pm	Software Setup Session	Edgar Lobaton


Introduction to Artificial Intelligence

Edgar Lobaton, Ph.D.

Director of Applied AI, College of Engineering (COE), NC State
Professor, Electrical and Computer (ECE) Engineering, NC State
NAIRR AI Education Fellow, Computing Research Association (CRA)


What is AI?

- AI allows computers to perform tasks that typically require human intelligence. AI systems can "learn" from information, recognize patterns, predict, make decisions and act on them.
- Examples:** Generative AI (ChatGPT, Meshy AI), recommendation (Netflix, Amazon), and code assistants (Codex, Claude Code).
- A course on AI studies and designs intelligent agents. These agents can perceive their environment, reason about information, and take autonomous actions to achieve a specific goal.




What is ML?

- **Machine Learning (ML)** is a subset of AI that focuses on algorithms that improve through experience. Instead of a human writing specific rules for every scenario, the programmer writes algorithms that allows the computer to find patterns in data.
- **Classical ML** requires feature engineering (i.e., humans tell the computer what parts of the data are important).
- **Deep Learning (DL)** is a subset of ML that discovers the features. Neural Networks are the foundation of DL.




Types of AI

- **Predictive AI: What will happen?**
 - Produces a forecast, score or classification
 - Uses pattern recognition on historical data
 - Examples: Predictive maintenance / weather forecast
- **Generative AI: What can be created?**
 - Produces new content (text, images, code)
 - Uses sampling from a learned distribution
 - Examples: Generate a design / draft a report



- **Agentic AI: What should be done?**
 - Completes multi-step tasks
 - Uses reasoning and tools through multiple iterations often taking advantage of generative models.
 - Examples: An agent that runs a simulation, checks results and revises without a human in the loop.

Agentic AI Pipeline




Types of AI - Scenarios

<https://forms.gle/4bUMCWrYhGrao3zCA>



Types of ML

- **Supervised Learning:**
What is the output for this input?
 - Uses labeled images
 - Learns a mapping from input to output
 - Example: Crack detection for labeled images
- **Unsupervised Learning:**
What pattern exists in this data?
 - Uses unlabeled examples
 - Finds groups, structure or distributions
 - Example: Clustering pump failures signatures from sensor logs



- **Reinforcement Learning:**
What action maximizes reward over time?
 - Learns through trial-and-error
 - Takes actions, receives rewards, adjusts behavior
 - Example: A robotic arm learning to grasp an object

Types of ML - Scenarios

<https://forms.gle/TaU3qWCbmXPi9ppg6>




Supervised Learning

Given that data is provided in input/output pairs, then ML can help identify the mapping between them. That is, $\text{output} = f(\text{input})$.




→ Cat

90 bpm




→ Dog

Stress




Supervised Learning



Think of a supervised learning model as a function or a black box (at least for now) that maps inputs (images) to outputs (labels in this case)




Artificial Neural Network (NN)




■ Single Neuron




Artificial Neural Network (NN)




Single Neuron



1-layer NN




K-layer NN




Deep Neural Network (DNN)

- A **DNN** is a NN with multiple layers and capable of learning features from raw data
- Traditional ML requires hand-craft features (e.g., computing statistical summaries); deep learning learns features directly from raw data, reducing that manual step
- Same core math as the 1980s backprop era — what changed is data availability, GPU compute, and architectural refinements (e.g., 2017 transformer architecture)



Understanding the Input and Output

Images are arrays of pixels value. All that the computer “sees” is a list of those numbers.



The labels are often encoded using a list as well. For example, if our list contains cats, dogs and birds, we could use a one-hot encoding:

$Cat \rightarrow [1 \ 0 \ 0]$; $Dog \rightarrow [0 \ 1 \ 0]$; $Bird \rightarrow [0 \ 0 \ 1]$

How Computers Learn

- Consider the training data
 - $[1 \ 1] \rightarrow 2$
 - $[2 \ 1] \rightarrow 3$
 - $[2 \ 2] \rightarrow 4$
 - What is the output for an input $[2 \ 1]$?
 - What is the output for an input $[3 \ 2]$?
- Consider the training
 - $[1 \ 1] \rightarrow 3$
 - $[2 \ 1] \rightarrow 5$
 - $[2 \ 2] \rightarrow 6$
 - What is the output for an input $[3 \ 2]$?
 - Relationship 1: $[x \ y] \rightarrow 2x + y = 8$
 - Relationship 2: $[x \ y] \rightarrow x^2 - x + y + 2 = 10$

We test an ML model by giving it data that it has not seen to ensure that it is capturing the underlying relationship. **[Generalization]**

The model needs to be properly specified given the amount of data provided. **[Overfitting / Underfitting]**

Let's Recap

Exploring a Code Assistant

Data: <https://drive.google.com/file/d/1RLQihFqsH8pfdVBdv7XGfpJzhP6ok0RO/view?usp=sharing>



Notebook: https://colab.research.google.com/drive/1loR3LFBVX3OnK8-1d3_IGwkFismnqg5y?usp=sharing

Join Discord

Data: <https://discord.gg/mJTMYWf6r>



Additional Info:

<https://docs.google.com/document/d/1pmE6s7hGsemdivZcKLxe7xCrPAY6hucoBYNCEXdqrO7k/edit?usp=sharing>

Submit your GitHub.com Account

Link: <https://forms.gle/Qpnh9SnhjXM4awLT7>

